



Neural Networks and Deep Learning

Lecture *Machine Learning* on March 9-10, 2026

Mario Stanke
Institut für Mathematik und Informatik
Universität Greifswald



Image classification problem

Given an image - say 64×64 pixels with 3 color channels - predict a probability distribution over predefined classes e.g. *cat*, *dog*, *cow*, *plant*, One can then assign the image to the most likely class.



Image classification problem

Given an image - say 64×64 pixels with 3 color channels - predict a probability distribution over predefined classes e.g. *cat*, *dog*, *cow*, *plant*, One can then assign the image to the most likely class.

Related problems

- Object detection
- Image segmentation
- Videos (i.e. timeseries of images) as input



Problems with Fully Connected Artificial Neural Nets (only Linear/Dense layers)

- high number of parameters



Problems with Fully Connected Artificial Neural Nets (only Linear/Dense layers)

- high number of parameters
- when images are input:



Problems with Fully Connected Artificial Neural Nets (only Linear/Dense layers)

- high number of parameters
- when images are input:
 - no notion of pixel neighborhoods
 - no translation invariance

Cross-correlation (2-dimensional)

Definition 1

Let $A = (a_{ij})_{0 \leq i, j < m}$ be a square $m \times m$ -dimensional matrix and

$$B = (b_{ij})_{\substack{0 \leq i < h \\ 0 \leq j < w}}$$

be another matrix of shape $h \times w$.

The $h - m + 1 \times w - m + 1$ -dimensional matrix C with entries

$$c_{i,j} := \sum_{i'=0}^{m-1} \sum_{j'=0}^{m-1} a_{i',j'} \cdot b_{i+i',j+j'}$$

is the 2-dimensional **cross-correlation** of A and B . We write $C = A * B$.

Cross-correlation (2-dimensional)

Definition 1

Let $A = (a_{ij})_{0 \leq i, j < m}$ be a square $m \times m$ -dimensional matrix and

$$B = (b_{ij})_{\substack{0 \leq i < h \\ 0 \leq j < w}}$$

be another matrix of shape $h \times w$.

The $h - m + 1 \times w - m + 1$ -dimensional matrix C with entries

$$c_{i,j} := \sum_{i'=0}^{m-1} \sum_{j'=0}^{m-1} a_{i',j'} \cdot b_{i+i',j+j'}$$

is the 2-dimensional **cross-correlation** of A and B . We write $C = A * B$.

Example 2


$m = 2$, $h = 4$, $w = 5$.

$$A = \begin{pmatrix} 1 & -1 \\ 2 & -3 \end{pmatrix} \quad B = \begin{pmatrix} 2 & -3 & 0 & 2 & -1 \\ 0 & 1 & 4 & 0 & 1 \\ 2 & -2 & 7 & 3 & 0 \\ -1 & 0 & 1 & 0 & 4 \end{pmatrix} \quad C = \begin{pmatrix} 2 & -13 & 6 & 0 \\ 9 & -28 & 9 & 5 \\ 2 & -12 & 6 & -9 \end{pmatrix}$$


Cross-Correlation of an Image

$$\begin{pmatrix} 0.56 & 0.54 & 0.59 & 0.41 & 0.56 & 0.54 & 0.45 & 0.50 & 0.44 & 0.58 & 0.44 & 0.54 \\ 0.48 & 0.42 & 0.53 & 0.48 & 0.53 & 0.46 & 0.42 & 0.52 & 0.60 & 0.52 & 0.56 & 0.52 \\ 0.40 & 0.56 & 0.56 & 0.42 & 0.51 & 0.55 & 0.58 & 0.47 & 0.55 & 0.53 & 0.51 & 0.57 \\ 0.40 & 0.44 & 0.58 & 1.00 & 0.51 & 0.47 & 0.51 & 0.43 & 0.54 & 0.48 & 0.56 & 0.46 \\ 0.53 & 0.42 & 0.58 & 1.00 & 0.59 & 0.55 & 0.56 & 0.43 & 0.57 & 0.52 & 0.53 & 0.52 \\ 0.44 & 0.53 & 0.45 & 1.00 & 0.51 & 0.43 & 0.43 & 0.53 & 0.49 & 0.42 & 0.52 & 0.57 \\ 0.41 & 0.51 & 0.53 & 1.00 & 0.46 & 0.43 & 0.54 & 0.42 & 0.57 & 0.42 & 0.53 & 0.53 \\ 0.43 & 0.49 & 0.49 & 1.00 & 0.48 & 0.46 & 0.56 & 0.48 & 0.50 & 0.54 & 0.56 & 0.57 \\ 0.45 & 0.49 & 0.44 & 1.00 & 1.00 & 1.00 & 1.00 & 1.00 & 1.00 & 0.46 & 0.59 & 0.58 \\ 0.59 & 0.45 & 0.50 & 0.48 & 0.48 & 0.60 & 0.46 & 0.41 & 0.46 & 0.55 & 0.57 & 0.40 \\ 0.56 & 0.48 & 0.45 & 0.57 & 0.55 & 0.49 & 0.48 & 0.45 & 0.47 & 0.50 & 0.58 & 0.43 \\ 0.41 & 0.48 & 0.44 & 0.54 & 0.43 & 0.55 & 0.52 & 0.54 & 0.55 & 0.43 & 0.53 & 0.59 \end{pmatrix}$$

“filter” A



input image B



output image



$$C = A * B$$


$$\begin{pmatrix} -1 & 4 & -1 & -1 & -1 \\ -1 & 4 & -1 & -1 & -1 \\ -1 & 4 & -1 & -1 & -1 \\ -1 & 4 & 4 & 4 & 4 \\ -1 & -1 & -1 & -1 & -1 \end{pmatrix}$$

Cross-Correlation of an Image

$$\begin{pmatrix} 0.56 & 0.54 & 0.59 & 0.41 & 0.56 & 0.54 & 0.45 & 0.50 & 0.44 & 0.58 & 0.44 & 0.54 \\ 0.48 & 0.42 & 0.53 & 0.48 & 0.53 & 0.46 & 0.42 & 0.52 & 0.60 & 0.52 & 0.56 & 0.52 \\ 0.40 & 0.56 & 0.56 & 0.42 & 0.51 & 0.55 & 0.58 & 0.47 & 0.55 & 0.53 & 0.51 & 0.57 \\ 0.40 & 0.44 & 0.58 & 1.00 & 0.51 & 0.47 & 0.51 & 0.43 & 0.54 & 0.48 & 0.56 & 0.46 \\ 0.53 & 0.42 & 0.58 & 1.00 & 0.59 & 0.55 & 0.56 & 0.43 & 0.57 & 0.52 & 0.53 & 0.52 \\ 0.44 & 0.53 & 0.45 & 1.00 & 0.51 & 0.43 & 0.43 & 0.53 & 0.49 & 0.42 & 0.52 & 0.57 \\ 0.41 & 0.51 & 0.53 & 1.00 & 0.46 & 0.43 & 0.54 & 0.42 & 0.57 & 0.42 & 0.53 & 0.53 \\ 0.43 & 0.49 & 0.49 & 1.00 & 0.48 & 0.46 & 0.56 & 0.48 & 0.50 & 0.54 & 0.56 & 0.57 \\ 0.45 & 0.49 & 0.44 & 1.00 & 1.00 & 1.00 & 1.00 & 1.00 & 1.00 & 0.46 & 0.59 & 0.58 \\ 0.59 & 0.45 & 0.50 & 0.48 & 0.48 & 0.60 & 0.46 & 0.41 & 0.46 & 0.55 & 0.57 & 0.40 \\ 0.56 & 0.48 & 0.45 & 0.57 & 0.55 & 0.49 & 0.48 & 0.45 & 0.47 & 0.50 & 0.58 & 0.43 \\ 0.41 & 0.48 & 0.44 & 0.54 & 0.43 & 0.55 & 0.52 & 0.54 & 0.55 & 0.43 & 0.53 & 0.59 \end{pmatrix}$$

“filter” A



$$\begin{pmatrix} -1 & 4 & -1 & -1 & -1 \\ -1 & 4 & -1 & -1 & -1 \\ -1 & 4 & -1 & -1 & -1 \\ -1 & 4 & 4 & 4 & 4 \\ -1 & -1 & -1 & -1 & -1 \end{pmatrix}$$



$C = A * B$

$$\begin{pmatrix} 6.6 & 7.4 & 5.0 & 4.19 & 4.7 & 4.3 & 4.7 & 5.2 \\ 6.2 & 7.9 & 8.8 & 4.56 & 5.3 & 5.4 & 4.6 & 6.2 \\ 5.2 & 5.9 & 9.2 & 3.27 & 4.6 & 5.2 & 3.8 & 5.6 \\ 4.7 & 5.2 & 12.1 & 2.57 & 4.6 & 5.0 & 4.0 & 5.5 \\ 4.4 & 4.1 & 11.0 & 0.85 & 2.0 & 3.7 & 3.1 & 5.1 \\ 7.7 & 9.3 & 19.3 & 11.13 & 11.9 & 10.8 & 8.4 & 7.2 \\ 2.7 & 2.7 & 9.6 & 3.69 & 4.4 & 5.7 & 5.4 & 6.7 \\ 3.7 & 3.1 & 7.9 & 4.14 & 4.8 & 5.2 & 5.3 & 6.0 \end{pmatrix}$$



3-dimensional input

- Want to
 - 1 use multiple filters in parallel and
 - 2 stack several (convolutional) layers.
- Also, color images are naturally encoded as 3-dimensional (each pixel has a red, green and blue value).
- Solution: Define convolution for 3-dimensional tensor input as well.






2-dimensional cross-correlation with channel dimension

Let $B = (b_{ijk})_{\substack{0 \leq i < h \\ 0 \leq j < w \\ 0 \leq k < d}}$ be a tensor of shape $h \times w \times d$ and

let $A = (a_{ijk})_{\substack{0 \leq i, j < m \\ 0 \leq k < d}}$ be another tensor ("filter").


The **cross-correlation** of A and B is then the $h - m + 1 \times w - m + 1$ -dimensional matrix $C = A * B$ with entries

$$c_{i,j} := \sum_{i'=0}^{m-1} \sum_{j'=0}^{m-1} \sum_{k=0}^{d-1} a_{i',j',k} \cdot b_{i+i',j+j',k}$$



Deep Learning for Computer Vision


Convolutional Layer



$$Z_{i,j} = \sum_{i'=0}^{m-1} \sum_{j'=0}^{m-1} \sum_{k=0}^{d-1} X_{i+i', j+j', k} \cdot f_{i', j', k}$$

Deep Learning for Computer Vision


Convolutional Layer




$$Z_{i,j,r} = \sum_{i'=0}^{m-1} \sum_{j'=0}^{m-1} \sum_{k=0}^{d-1} x_{i+i',j+j',k} \cdot f_{i',j',k,r} + \underset{\text{bias}}{b_r}$$

Deep Learning for Computer Vision


Convolutional Layer



Convolutional Layer




Convolutional Layer




- The input width and height can be conserved in the output layer by zero-padding of input (`padding = 'same'`)

Convolutional Layer




- The input width and height can be conserved in the output layer by zero-**padding** of input (`padding = 'same'`)
- **Stride** (Schrittweite) s : Skip $s - 1$ positions in each direction when 'sliding' A over $B \Rightarrow$ decreases output layer size up to a factor of s^2 .

Convolutional Layer




- The input width and height can be conserved in the output layer by zero-**padding** of input (`padding = 'same'`)
- **Stride** (Schrittweite) s : Skip $s - 1$ positions in each direction when 'sliding' A over $B \Rightarrow$ decreases output layer size up to a factor of s^2 .
- The matrices A are **learned**, not set manually. The derivative wrt. to the filter matrix parameters is computed by the backpropagation algorithm.

Convolutional Layer



- The input width and height can be conserved in the output layer by zero-**padding** of input (`padding = 'same'`)
- **Stride** (Schrittweite) s : Skip $s - 1$ positions in each direction when 'sliding' A over $B \Rightarrow$ decreases output layer size up to a factor of s^2 .
- The matrices A are **learned**, not set manually. The derivative wrt. to the filter matrix parameters is computed by the backpropagation algorithm.
- Convolution is a special case of a fully-connected layer, in which certain parameters are shared (*parameter sharing*).

Convolutional Layer



- The input width and height can be conserved in the output layer by zero-**padding** of input (`padding = 'same'`)
- **Stride** (Schrittweite) s : Skip $s - 1$ positions in each direction when 'sliding' A over $B \Rightarrow$ decreases output layer size up to a factor of s^2 .
- The matrices A are **learned**, not set manually. The derivative wrt. to the filter matrix parameters is computed by the backpropagation algorithm.
- Convolution is a special case of a fully-connected layer, in which certain parameters are shared (*parameter sharing*).
- Output neurons of convolution can detect lower-level features like ("lower left corner", "pupil") and be combined in deeper layers.

Convolutional Layer for Sequences (Conv1D, 1-dim. cross-correlation with channel dim.)

- Let's directly consider the important case `padding='same'`. Why may one want to preserve the sequence length exactly over the layers?
- For symmetry consider an *odd* kernel size $m = 2r + 1$

Let

$$B = (b_{jk}) \quad \begin{matrix} 0 \leq j < w \\ 0 \leq k < d \end{matrix}$$

be the input "sequence" and let

$$A = (a_{jk}) \quad \begin{matrix} 0 \leq j < m \\ 0 \leq k < d \end{matrix}$$

be a filter.

The **cross-correlation** of A and B is then the


$$w - m + 1$$

-dimensional vector $C = A * B$ with entries

$$c_j := \sum_{j'=0}^{m-1} \sum_{k=0}^{d-1} a_{j',k} \cdot b_{j+j'-r,k} \quad (0 \leq j < w)$$

With **stride** s and so-called **dilation** ℓ (s, ℓ in $\{1, 2, 3, \dots\}$) we compute

$$c_j := \sum_{j'=0}^{m-1} \sum_{k=0}^{d-1} a_{j',k} \cdot b_{sj+\ell(j'-r),k}$$





Pooling-Layers

Max-Pooling (MaxPool2D)

- similar to a convolutional layer
- requires a `pool_size` m like the filter size
- does not have any parameters
- computes output

$$Z_{i,j,r} = \max_{\substack{i' \in [0, m) \\ j' \in [0, m)}} X_{s \cdot i + i', s \cdot j + j', r}$$

- is usually applied with a **stride** $s \geq 2$ and therefore reduces height and width
- often $s = m$, can have different strides for each dimension
- intuition:



Pooling-Layers

Max-Pooling (MaxPool2D)

- similar to a convolutional layer
- requires a `pool_size` m like the filter size
- does not have any parameters
- computes output

$$Z_{i,j,r} = \max_{\substack{i' \in [0, m) \\ j' \in [0, m)}} X_{s \cdot i + i', s \cdot j + j', r}$$

- is usually applied with a **stride** $s \geq 2$ and therefore reduces height and width
- often $s = m$, can have different strides for each dimension
- intuition: checks local presence of *at least one* feature (e.g. a green stroke) rather than summing up if there are multiple



Pooling-Layers

Max-Pooling (MaxPool2D)

- similar to a convolutional layer
- requires a `pool_size` m like the filter size
- does not have any parameters
- computes output

$$Z_{i,j,r} = \max_{\substack{i' \in [0, m) \\ j' \in [0, m)}} X_{s \cdot i + i', s \cdot j + j', r}$$


- is usually applied with a **stride** $s \geq 2$ and therefore reduces height and width
- often $s = m$, can have different strides for each dimension
- intuition: checks local presence of *at least one* feature (e.g. a green stroke) rather than summing up if there are multiple

With an analogous definition, **average pooling** averages over regions of size $m \times m$, but is used less often.



Photo classification


- CNN from 2012 ("AlexNet")
- classification into 1000 categories



Alex Krizhevsky, Ilya Sutskever and Geoffrey Hinton, "ImageNet Classification with Deep Convolutional Neural Networks", NIPS, 2012

Image segmentation, diffusion model


- outputs images, e.g. the segmented input image
- part of stable diffusion text2image model



Olaf Ronneberger, Philipp Fischer, Thomas Brox



Stability AI



U-net: Convolutional networks for biomedical image segmentation

O Ronneberger, P Fischer, T Brox - Medical image computing and computer-assisted ..., 2015

Cited by 106928