



Machine Learning

Lecture *Machine Learning* on March 9-10, 2026

Artificial Neural Networks

Definition

Training (BackProp)

Activation Functions

Mario Stanke
Institut für Mathematik und Informatik
Universität Greifswald

Definition 1 (Feed-Forward Neural Network (MLP = Multilayer Perceptron))

A feed-forward **artificial neural network** with $L \geq 1$ layers of sizes $s_1, \dots, s_L$ with $n =: s_0$ input variables $\mathbf{x} = (x_1, \dots, x_n)^T$ and $K := s_L$ output variables $\mathbf{t} = (t_1, \dots, t_K)^T$ is a function

$$\mathbf{t} = h_{\boldsymbol{\theta}}(\mathbf{x})$$

with parameters


$$\boldsymbol{\theta} = (\Theta^{(0)}, \dots, \Theta^{(L-1)}) \quad , \text{ where } \Theta^{(\ell)} \in \mathbb{R}^{s_{\ell+1} \times (s_{\ell}+1)},$$

defined by the following recursions

$$\begin{aligned} \mathbf{t} &= \mathbf{g}(\mathbf{z}^{(L)}) \in \mathbb{R}^K \\ \mathbf{z}^{(\ell)} &= \Theta^{(\ell-1)} \mathbf{a}^{(\ell-1)} \in \mathbb{R}^{s_{\ell}} \quad (1 \leq \ell \leq L) \\ \mathbf{a}^{(\ell)} &= \begin{pmatrix} a_0^{(\ell)} \\ a_1^{(\ell)} \\ \vdots \\ a_{s_{\ell}}^{(\ell)} \end{pmatrix} = \begin{pmatrix} 1 \\ \sigma(z_1^{(\ell)}) \\ \vdots \\ \sigma(z_{s_{\ell}}^{(\ell)}) \end{pmatrix} \quad (1 \leq \ell < L) \\ \mathbf{a}^{(0)} &= \begin{pmatrix} 1 \\ x_1 \\ \vdots \\ x_n \end{pmatrix} . \end{aligned}$$

Here, σ is called an **activation function** and we will call $\mathbf{g} : \mathbb{R}^K \rightarrow \mathbb{R}^K$ the **output activation function**.

Linear regression = MLP with 1 layer and 1 output variable



$L = 1$ layers

$\theta = (\theta_0, \theta_1, \dots, \theta_n)$ parameters


t_1 one output unit

$x_1, \dots, x_n$ input units

output activation function $t_1 = g(z_1) = z_1$

$$t_1 = \theta_0 + \sum_{j=1}^n \theta_j x_j$$

Linear regression = MLP with 1 layer and 1 output variable



$L = 1$ layers

$\theta = (\theta_0, \theta_1, \dots, \theta_n)$ parameters

t_1 one output unit


$x_1, \dots, x_n$ input units

output activation function $t_1 = g(z_1) = z_1$

$$t_1 = \theta_0 + \sum_{j=1}^n \theta_j x_j$$

Artificial Neural Networks

MLP with 1 hidden layer



$L = 2$ layers, 1 hidden layer, $s_1 = 2$


$\theta = (\Theta^{(0)}, \Theta^{(1)})$ parameters

x_1, x_2 input units, t_1 single output unit with output activation function σ

$$t_1 = \sigma \left(\theta_{10}^{(1)} + \theta_{11}^{(1)} a_1 + \theta_{12}^{(1)} a_2 \right) = \sigma \left(\theta_{10}^{(1)} + \theta_{11}^{(1)} \sigma(\theta_{10}^{(0)} + \theta_{11}^{(0)} x_1 + \theta_{12}^{(0)} x_2) + \theta_{12}^{(1)} \sigma(\theta_{20}^{(0)} + \theta_{21}^{(0)} x_1 + \theta_{22}^{(0)} x_2) \right)$$

Artificial Neural Networks

MLP with 1 hidden layer



$L = 2$ layers, 1 hidden layer, $s_1 = 2$


$\theta = (\Theta^{(0)}, \Theta^{(1)})$ parameters

x_1, x_2 input units, t_1 single output unit with output activation function σ

$$t_1 = \sigma \left(\theta_{10}^{(1)} + \theta_{11}^{(1)} a_1 + \theta_{12}^{(1)} a_2 \right) = \sigma \left(\theta_{10}^{(1)} + \theta_{11}^{(1)} \sigma(\theta_{10}^{(0)} + \theta_{11}^{(0)} x_1 + \theta_{12}^{(0)} x_2) + \theta_{12}^{(1)} \sigma(\theta_{20}^{(0)} + \theta_{21}^{(0)} x_1 + \theta_{22}^{(0)} x_2) \right)$$

Artificial Neural Networks

MLP with 1 hidden layer



$L = 2$ layers, 1 hidden layer, $s_1 = 2$

$\theta = (\Theta^{(0)}, \Theta^{(1)})$ parameters

x_1, x_2 input units, t_1 single output unit with output activation function σ

$$t_1 = \sigma \left(\theta_{10}^{(1)} + \theta_{11}^{(1)} a_1 + \theta_{12}^{(1)} a_2 \right) = \sigma \left(\theta_{10}^{(1)} + \theta_{11}^{(1)} \sigma(\theta_{10}^{(0)} + \theta_{11}^{(0)} x_1 + \theta_{12}^{(0)} x_2) + \theta_{12}^{(1)} \sigma(\theta_{20}^{(0)} + \theta_{21}^{(0)} x_1 + \theta_{22}^{(0)} x_2) \right)$$

Training of Artificial Neural Networks

Suppose a *single* training observation

$$\mathbf{y} = (y_1, \dots, y_K)^T$$

for training input

$$\mathbf{x} = (x_1, \dots, x_n)^T$$

is given.

Training of Artificial Neural Networks

Suppose a *single* training observation

$$\mathbf{y} = (y_1, \dots, y_K)^T$$

for training input

$$\mathbf{x} = (x_1, \dots, x_n)^T$$

is given.

Consider some error function

$$D(\theta)$$

that typically depends on training data but that depends on the parameters θ only through the network output $\mathbf{t}$, e.g. a squared error function.

Training of Artificial Neural Networks

Suppose a *single* training observation

$$\mathbf{y} = (y_1, \dots, y_K)^T$$

for training input

$$\mathbf{x} = (x_1, \dots, x_n)^T$$

is given.

Consider some error function

$$D(\boldsymbol{\theta})$$

that typically depends on training data but that depends on the parameters $\boldsymbol{\theta}$ only through the network output $\mathbf{t}$, e.g. a squared error function. To minimize the $D(\boldsymbol{\theta})$ we will require to compute all partial derivatives

$$\frac{\partial D}{\partial \theta_{ji}^{(\ell-1)}}$$

for ℓ , j and i (“gradient” in TensorFlow).

Training of Artificial Neural Networks

Suppose a *single* training observation

$$\mathbf{y} = (y_1, \dots, y_K)^T$$

for training input

$$\mathbf{x} = (x_1, \dots, x_n)^T$$

is given.

Consider some error function

$$D(\boldsymbol{\theta})$$


that typically depends on training data but that depends on the parameters $\boldsymbol{\theta}$ only through the network output $\mathbf{t}$, e.g. a squared error function. To minimize the $D(\boldsymbol{\theta})$ we will require to compute all partial derivatives

$$\frac{\partial D}{\partial \theta_{ji}^{(\ell-1)}}$$

for ℓ , j and i (“gradient” in TensorFlow).

$D(\boldsymbol{\theta})$ may be non-convex and have **local, non-global minima**.

NN Training: partial derivatives **backpropagate** through network



From

$$z_j^{(\ell)} = \sum_r \theta_{jr}^{(\ell-1)} a_r^{(\ell-1)}$$

and using the multivariate chain rule that

we obtain using the univariate chain rule that

$$\frac{\partial D}{\partial \theta_{ji}^{(\ell-1)}} = \frac{\partial D}{\partial z_j^{(\ell)}} a_i^{(\ell-1)}$$

$$\frac{\partial D}{\partial z_j^{(\ell)}} = \left(\frac{\partial D}{\partial z_1^{(\ell+1)}}, \dots, \frac{\partial D}{\partial z_{s_{\ell+1}}^{(\ell+1)}} \right) \begin{pmatrix} \frac{\partial z_1^{(\ell+1)}}{\partial z_j^{(\ell)}} \\ \vdots \\ \frac{\partial z_{s_{\ell+1}}^{(\ell+1)}}{\partial z_j^{(\ell)}} \end{pmatrix}$$

All derivatives can be computed efficiently in one right-to-left pass (“**backpropagation** algorithm”). TensorFlow does backprop automatically, for general models.

Artificial Neural
Networks

Definition

Training (BackProp)

Activation Functions

Artificial Neural Networks for Regression

Consider first a **single training instance** $\mathbf{x} \in R^n$ with a single output $\mathbf{y} \in R^K$.

(We simply average $D(\theta)$ over multiple training instances.)



Artificial Neural Networks for Regression

Consider first a **single training instance** $\mathbf{x} \in \mathbb{R}^n$ with a single output $\mathbf{y} \in \mathbb{R}^K$.

(We simply average $D(\theta)$ over multiple training instances.)

- Choose the **identity function** as output activation function $\mathbf{g}$, i.e. $\mathbf{t} = \mathbf{z}^{(L)}$ (unbounded).



Artificial Neural Networks for Regression

Consider first a **single training instance** $\mathbf{x} \in R^n$ with a single output $\mathbf{y} \in R^K$.

(We simply average $D(\theta)$ over multiple training instances.)

- Choose the **identity function** as output activation function $\mathbf{g}$, i.e. $\mathbf{t} = \mathbf{z}^{(L)}$ (unbounded).
- Choose the **squared error function**

$$D(\theta) := \|\mathbf{t} - \mathbf{y}\|_2^2 = \sum_{k=1}^K (t_k - y_k)^2$$

Artificial Neural Networks for Classification

Again, consider first a single training instance. Here, let output $\mathbf{y} \in \mathbb{R}^K$ be one-hot encoded such that

$$\mathbf{y} = \mathbf{e}_c = c\text{-th unit vector}$$

if $c \in \{1, 2, \dots, K\}$ is the true class of the learning instance.

Artificial Neural Networks for Classification

Again, consider first a single training instance. Here, let output $\mathbf{y} \in \mathbb{R}^K$ be **one-hot** encoded such that

$$\mathbf{y} = \mathbf{e}_c = c\text{-th unit vector}$$

if $c \in \{1, 2, \dots, K\}$ is the true class of the learning instance.

- Choose the **softmax** function as output activation function $\mathbf{g}$.

$$\mathbf{t} = \mathbf{g}(z_1, \dots, z_K) = \frac{1}{\sum_{k=1}^K e^{z_k}} \begin{pmatrix} e^{z_1} \\ \vdots \\ e^{z_K} \end{pmatrix}$$

Artificial Neural Networks for Classification

Again, consider first a single training instance. Here, let output $\mathbf{y} \in R^K$ be **one-hot** encoded such that

$$\mathbf{y} = \mathbf{e}_c = c\text{-th unit vector}$$

if $c \in \{1, 2, \dots, K\}$ is the true class of the learning instance.

- Choose the **softmax** function as output activation function $\mathbf{g}$.

$$\mathbf{t} = \mathbf{g}(z_1, \dots, z_K) = \frac{1}{\sum_{k=1}^K e^{z_k}} \begin{pmatrix} e^{z_1} \\ \vdots \\ e^{z_K} \end{pmatrix}$$

Properties:

- $0 < t_k < 1$, $t_1 + \dots + t_K = 1$
- if $z_j \gg z_i$ for $i \neq j$, then $t_j \approx 1$ and $t_i \approx 0$ for $i \neq j$
- softmax is a generalization of the sigmoid function to more than 2 classes
- If $K = 2$ (binary classification), $t_1 = 1 - t_2$ is often not stored. The same neural network function can then be achieved by setting the output size to 1 and using the logistic sigmoid function as activation.

Artificial Neural Networks for Classification

Cross-entropy error function for multiclass problem

Consider the true class of the training instance to be a random variable $C \in \{1, \dots, K\}$ with observation c and the output vector $\mathbf{t}$ of the net to be the parameters of a multinomial distribution for C . Then the likelihood is

$$L(\theta) =$$

Artificial Neural Networks for Classification

Cross-entropy error function for multiclass problem

Consider the true class of the training instance to be a random variable $C \in \{1, \dots, K\}$ with observation c and the output vector $\mathbf{t}$ of the net to be the parameters of a multinomial distribution for C . Then the likelihood is

$$L(\theta) = t_c$$

Artificial Neural Networks for Classification

Cross-entropy error function for multiclass problem

Consider the true class of the training instance to be a random variable $C \in \{1, \dots, K\}$ with observation c and the output vector $\mathbf{t}$ of the net to be the parameters of a multinomial distribution for C . Then the likelihood is

$$L(\boldsymbol{\theta}) = t_c$$

and we seek to minimize the negative log-likelihood

$$-\ln L(\boldsymbol{\theta}) = -\ln t_c = -\sum_{k=1}^K y_k \ln t_k$$

Artificial Neural Networks for Classification

Cross-entropy error function for multiclass problem

Consider the true class of the training instance to be a random variable $C \in \{1, \dots, K\}$ with observation c and the output vector $\mathbf{t}$ of the net to be the parameters of a multinomial distribution for C . Then the likelihood is

$$L(\boldsymbol{\theta}) = t_c$$

and we seek to minimize the negative log-likelihood

$$-\ln L(\boldsymbol{\theta}) = -\ln t_c = -\sum_{k=1}^K y_k \ln t_k$$

Define the **cross-entropy** error function

$$D(\boldsymbol{\theta}) = -\sum_{k=1}^K y_k \ln t_k.$$



Artificial Neural Networks

Training set

Let

$$\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}$$

be training **inputs** and

$$\mathbf{y}^{(1)}, \dots, \mathbf{y}^{(m)}$$

be the corresponding training **outputs / labels**. Let

$$\mathbf{t}^{(i)} = h_{\Theta}(\mathbf{x}^{(i)}) \quad (i = 1..m)$$

be the outputs of the MLP.

Overfitting

If the model has many parameters, the training may result in a model that fits the training data 'too well', therefore **does not generalize well** and performs poorly on **independent** test data. Remedies:

- ① Make model less complex, e.g. reduce number of parameters or change model class (e.g. lin. regresion over NN).
- ② **Regularize** the model: Penalize certain parameter values independent of the data.

Overfitting

If the model has many parameters, the training may result in a model that fits the training data 'too well', therefore **does not generalize well** and performs poorly on **independent** test data. Remedies:

- ① Make model less complex, e.g. reduce number of parameters or change model class (e.g. lin. regression over NN).
- ② **Regularize** the model: Penalize certain parameter values independent of the data.

Error function with regularization for neural networks

$$E(\theta; X) = D(\theta; X) + \lambda R(\theta)$$

Overfitting

If the model has many parameters, the training may result in a model that fits the training data 'too well', therefore **does not generalize well** and performs poorly on **independent** test data. Remedies:

- 1 Make model less complex, e.g. reduce number of parameters or change model class (e.g. lin. regression over NN).
- 2 **Regularize** the model: Penalize certain parameter values independent of the data.

Error function with regularization for neural networks

$$E(\theta; X) = D(\theta; X) + \lambda R(\theta)$$

An common choice for the regularization term is the (scaled) L2 norm:

$$R(\theta) = \frac{1}{m} \sum_{\ell=0}^{L-1} \sum_{i=1}^{s_{\ell}} \sum_{j=1}^{s_{\ell+1}} \left(\theta_{ji}^{(\ell)} \right)^2.$$

Note that the bias terms $\theta_{j0}^{(\ell)}$ are left out (not penalized).

Artificial Neural Networks

Error function with regularization for regression

$$E(\theta) = \frac{1}{m} \left\{ \sum_{i=1}^m \sum_{k=1}^K \left(t_k^{(i)} - y_k^{(i)} \right)^2 + \lambda \sum_{\ell=0}^{L-1} \sum_{i=1}^{s_{\ell}} \sum_{j=1}^{s_{\ell+1}} \left(\theta_{ji}^{(\ell)} \right)^2 \right\} \quad (1)$$

Artificial Neural Networks

Error function with regularization for regression

$$E(\theta) = \frac{1}{m} \left\{ \sum_{i=1}^m \sum_{k=1}^K \left(t_k^{(i)} - y_k^{(i)} \right)^2 + \lambda \sum_{\ell=0}^{L-1} \sum_{i=1}^{s_{\ell}} \sum_{j=1}^{s_{\ell+1}} \left(\theta_{ji}^{(\ell)} \right)^2 \right\} \quad (1)$$

Error function with regularization for classification

$$E(\theta) = \frac{1}{m} \left\{ - \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \ln t_k^{(i)} + \lambda \sum_{\ell=0}^{L-1} \sum_{i=1}^{s_{\ell}} \sum_{j=1}^{s_{\ell+1}} \left(\theta_{ji}^{(\ell)} \right)^2 \right\} \quad (2)$$




Artificial Neural Networks

Definition

Training (BackProp)

Activation Functions

1 hidden layer of size 5, $\lambda = 0$ (no regularization)




Artificial Neural Networks

Definition

Training (BackProp)

Activation Functions

2 hidden layers of size 8 each, $\lambda = 0$ (no regularization)




Artificial Neural Networks

Definition

Training (BackProp)

Activation Functions

2 hidden layers of size 20 each, $\lambda = 0$ (no regularization)




Artificial Neural Networks

Definition

Training (BackProp)

Activation Functions

2 hidden layers of size 20 each, $\lambda = 0.1$ (with regularization)




Artificial Neural Networks

Definition

Training (BackProp)

Activation Functions

Posterior probability of class (theoretical optimum if distr. was known)





Artificial Neural Networks

Definition

Training (BackProp)


Activation Functions

Activation Functions

Logistic sigmoid, `tf.nn.sigmoid`

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

$$\sigma' = \sigma(1 - \sigma)$$




Tangens hyperbolicus

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

$$= 2\sigma(2x) - 1$$

$$\tanh'(x) = 1 - (\tanh(x))^2$$





Artificial Neural Networks

Definition

Training (BackProp)


Activation Functions

Rectified linear unit (ReLU), `tf.nn.relu`

$$\text{relu}(x) = x^+ = \max\{0, x\}$$

$$\text{relu}'(x) = \begin{cases} 1 & , \text{if } x > 0 \\ 0 & , \text{if } x < 0 \end{cases}$$

$$\text{relu}'(0) := 1 \quad (\text{arbitrary})$$


Leaky rectified linear unit (Leaky ReLU), `tf.nn.leaky_relu`

$$\text{lrelu}(x) = \max\{\alpha x, x\}$$

$$\text{lrelu}'(x) = \begin{cases} 1 & , \text{if } x > 0 \\ \alpha & , \text{if } x < 0 \end{cases}$$

$$\text{lrelu}'(0) := 1 \quad (\text{arbitrary})$$

$$\alpha = 0.01$$



Artificial Neural
Networks

Definition

Training (BackProp)

Activation Functions

Activation Functions

Vanishing gradients

- vanishing gradient problem:
 - during backprop in each layer, a derivative of the activation function is multiplied
 - for NNs with many layers (deep networks) this can result in gradients that are practically zero (“vanishing”), in particular for the sigmoid function
- ReLU
 - most frequently chosen
 - introduced in 2010

(Nair and Hinton, “Rectified linear units improve restricted Boltzmann machines”, *Proc. ICML*, 2010)

- The derivative is either 0 or 1. Can lead to sparse gradients.